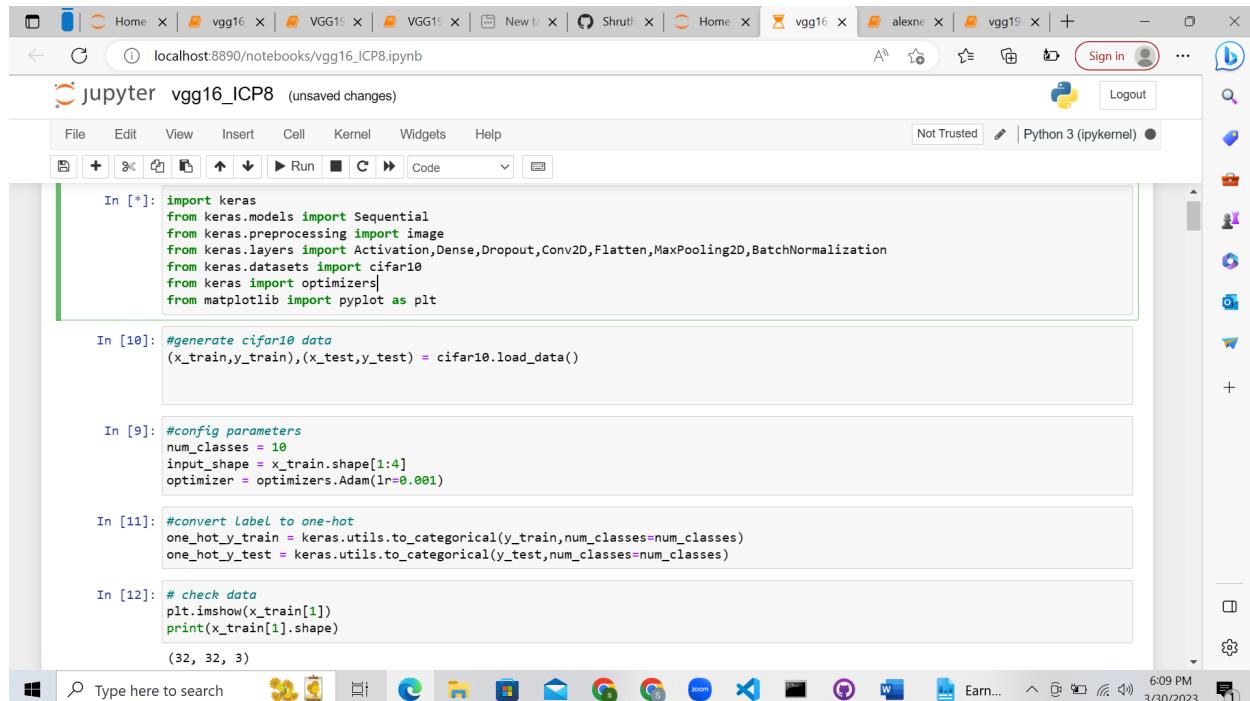


NN & DeepLearning : Image Classification with CNN_ICP8

In class programming:

1. Tune hyperparameter and make necessary addition to the baseline model to improve validation accuracy and reduce validation loss.
2. Provide logical description of which steps lead to improved response and what was its impact on architecture behavior.
3. Create at least two more visualizations using matplotlib (Other than provided in the source file)
4. Use dataset of your own choice and implement baseline models provided.
5. Apply modified architecture to your own selected dataset and train it.
6. Evaluate your model on testing set.
7. Save the improved model and use it for prediction on testing data
8. Provide plot of confusion matrix
9. Provide Training and testing Loss and accuracy plots in one plot using subplot command and history object.
10. Provide at least two more visualizations reflecting your solution.
11. Provide logical description of which steps lead to improved response for new dataset when compared with baseline model and enhance architecture and what was its impact on architecture behavior

Ans



The screenshot shows a Jupyter Notebook interface with the following code cells:

```
In [1]: import keras
       from keras.models import Sequential
       from keras.preprocessing import image
       from keras.layers import Activation,Dense,Dropout,Conv2D,Flatten,MaxPooling2D,BatchNormalization
       from keras.datasets import cifar10
       from keras import optimizers
       from matplotlib import pyplot as plt

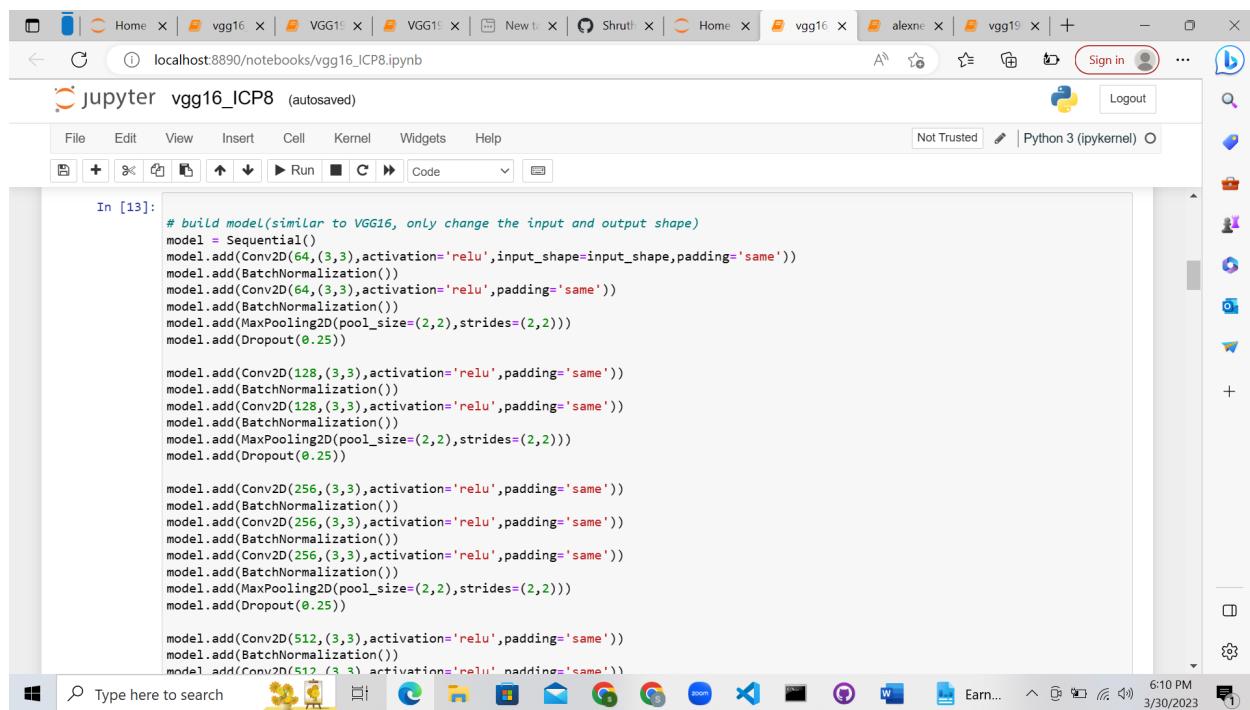
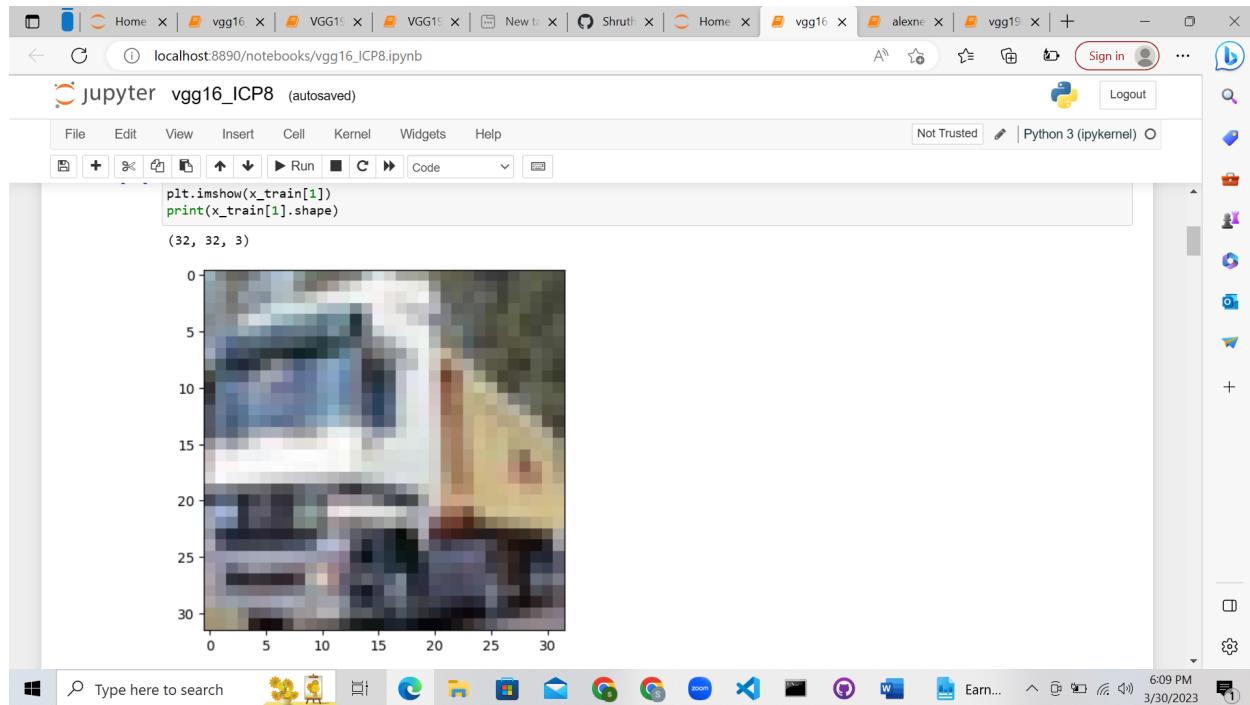
In [10]: #generate cifar10 data
        (x_train,y_train),(x_test,y_test) = cifar10.load_data()

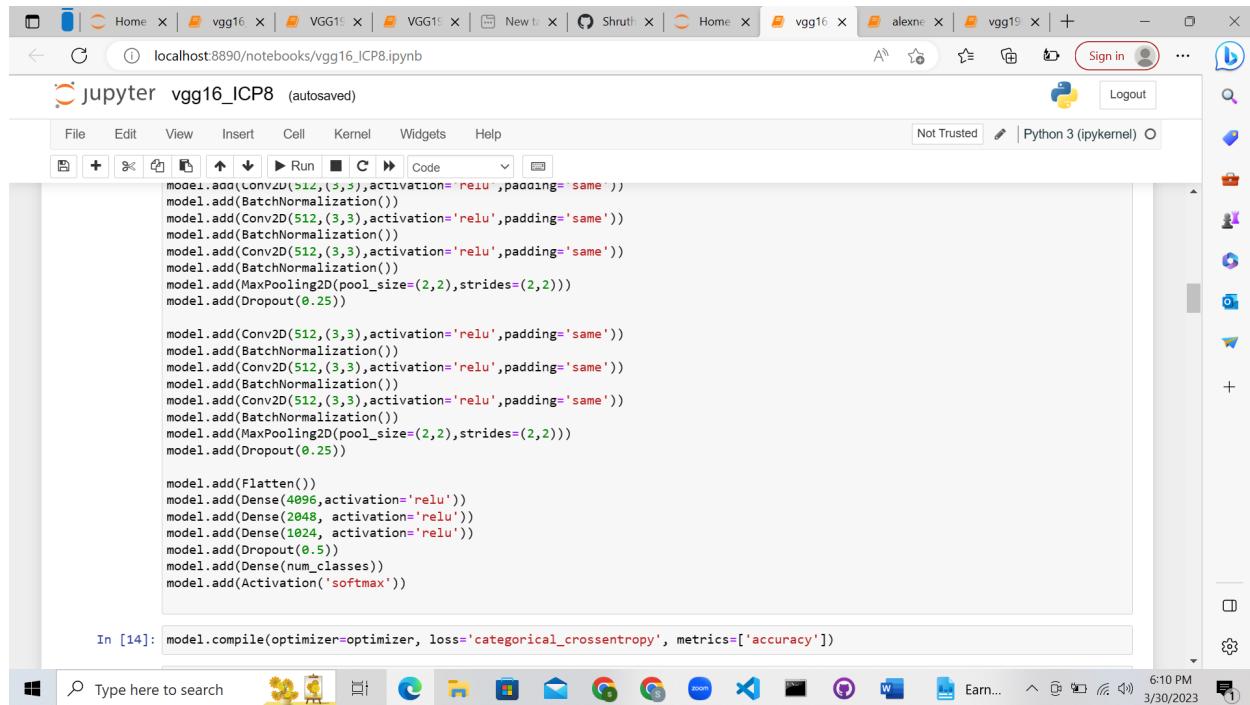
In [9]: #config parameters
        num_classes = 10
        input_shape = x_train.shape[1:4]
        optimizer = optimizers.Adam(lr=0.001)

In [11]: #convert label to one-hot
        one_hot_y_train = keras.utils.to_categorical(y_train,num_classes=num_classes)
        one_hot_y_test = keras.utils.to_categorical(y_test,num_classes=num_classes)

In [12]: # check data
        plt.imshow(x_train[1])
        print(x_train[1].shape)
        (32, 32, 3)
```

The notebook is running on a Python 3 kernel. The status bar at the bottom right shows the time as 6:09 PM and the date as 3/30/2023.



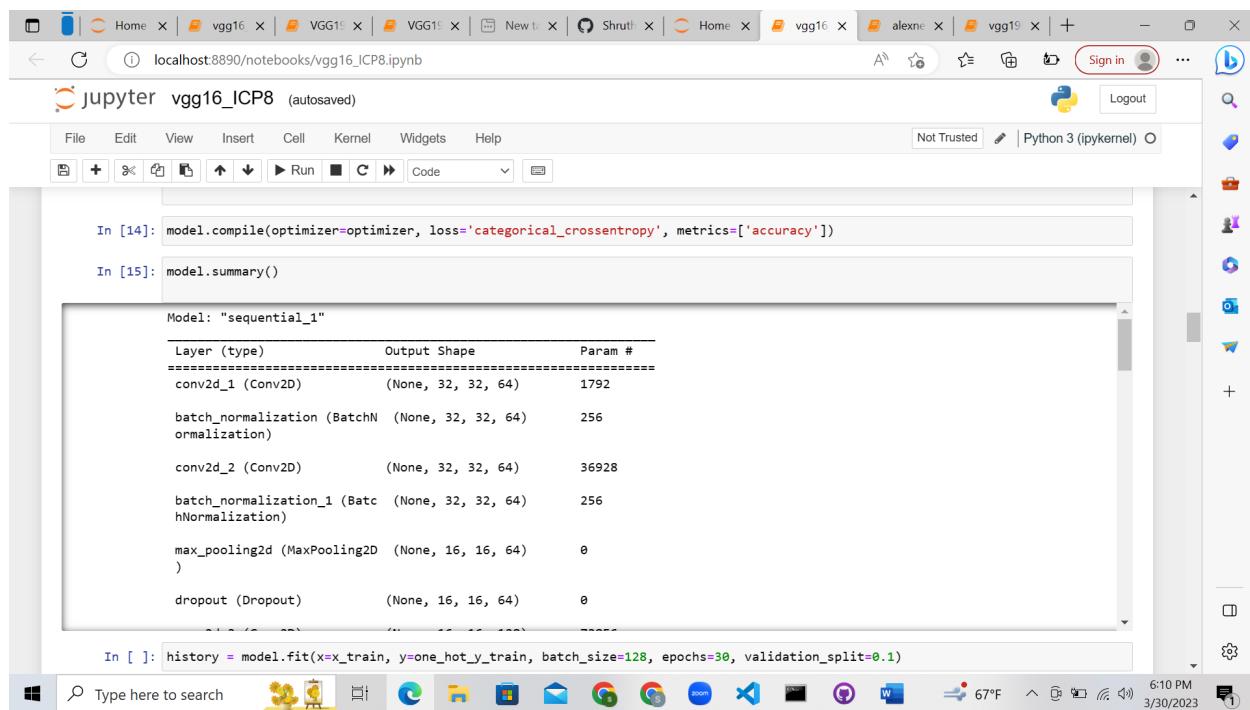


```
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(4096,activation='relu'))
model.add(Dense(2048,activation='relu'))
model.add(Dense(1024,activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

In [14]: model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
```



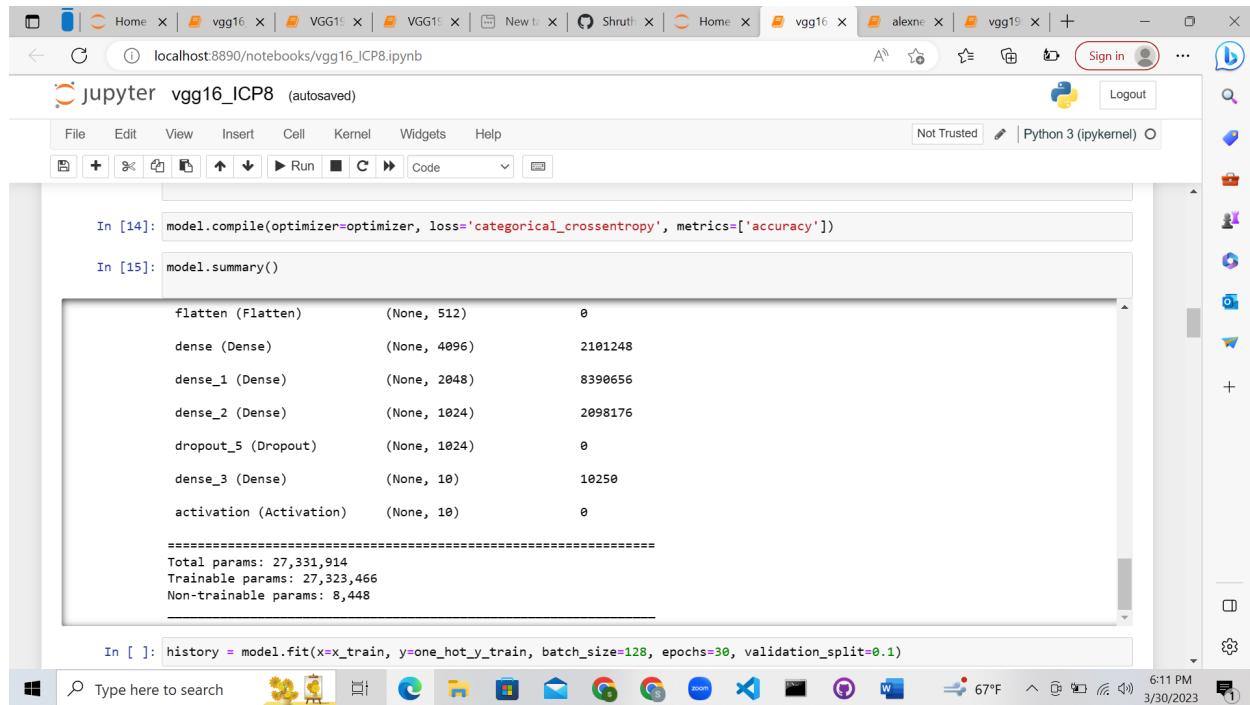
Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 64)	1792
batch_normalization (BatchN ormalization)	(None, 32, 32, 64)	256
conv2d_2 (Conv2D)	(None, 32, 32, 64)	36928
batch_normalization_1 (BatchN ormalization)	(None, 32, 32, 64)	256
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0
dropout (Dropout)	(None, 16, 16, 64)	0

```
In [14]: model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

In [15]: model.summary()

Model: "sequential_1"
Layer (type)          Output Shape         Param #
=====
conv2d_1 (Conv2D)     (None, 32, 32, 64)  1792
batch_normalization (BatchN ormalization) (None, 32, 32, 64)  256
conv2d_2 (Conv2D)     (None, 32, 32, 64)  36928
batch_normalization_1 (BatchN ormalization) (None, 32, 32, 64)  256
max_pooling2d (MaxPooling2D) (None, 16, 16, 64)  0
dropout (Dropout)     (None, 16, 16, 64)  0

In [16]: history = model.fit(x=x_train, y=one_hot_y_train, batch_size=128, epochs=30, validation_split=0.1)
```



In [14]: `model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])`

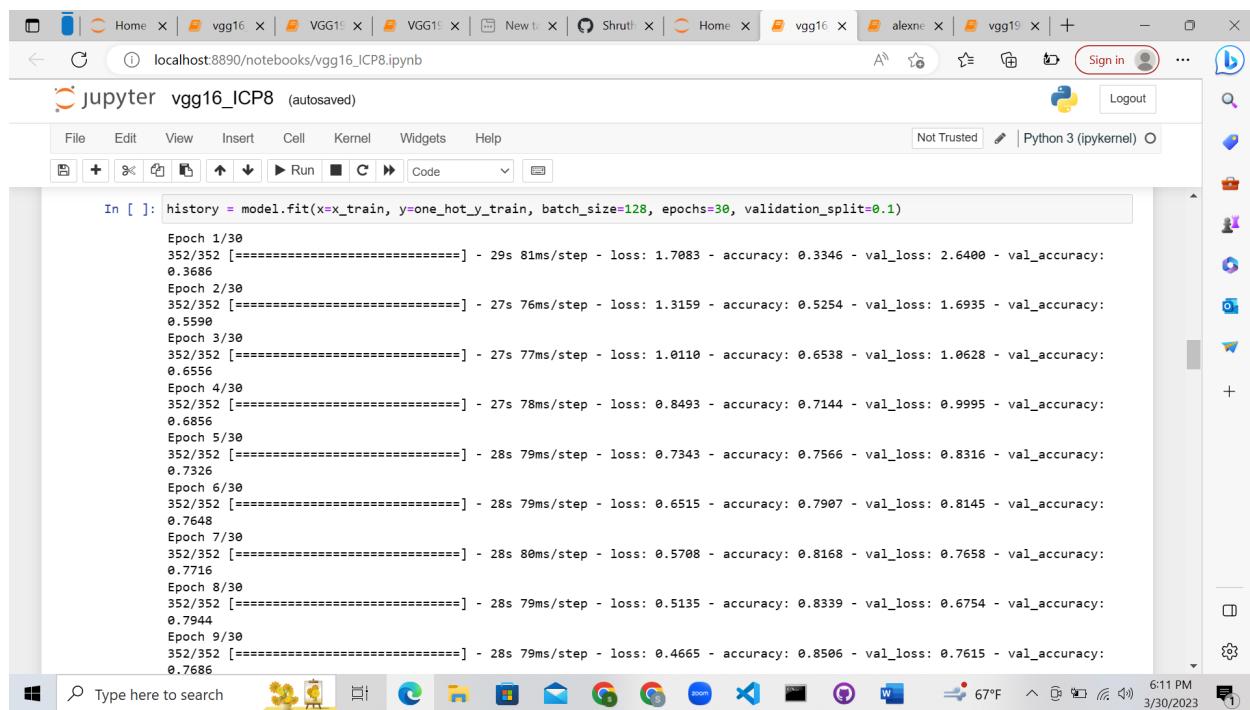
In [15]: `model.summary()`

Layer	Type	Shape	Count
flatten (Flatten)	(None, 512)	0	
dense (Dense)	(None, 4096)	2101248	
dense_1 (Dense)	(None, 2048)	8390656	
dense_2 (Dense)	(None, 1024)	2098176	
dropout_5 (Dropout)	(None, 1024)	0	
dense_3 (Dense)	(None, 10)	10250	
activation (Activation)	(None, 10)	0	

=====

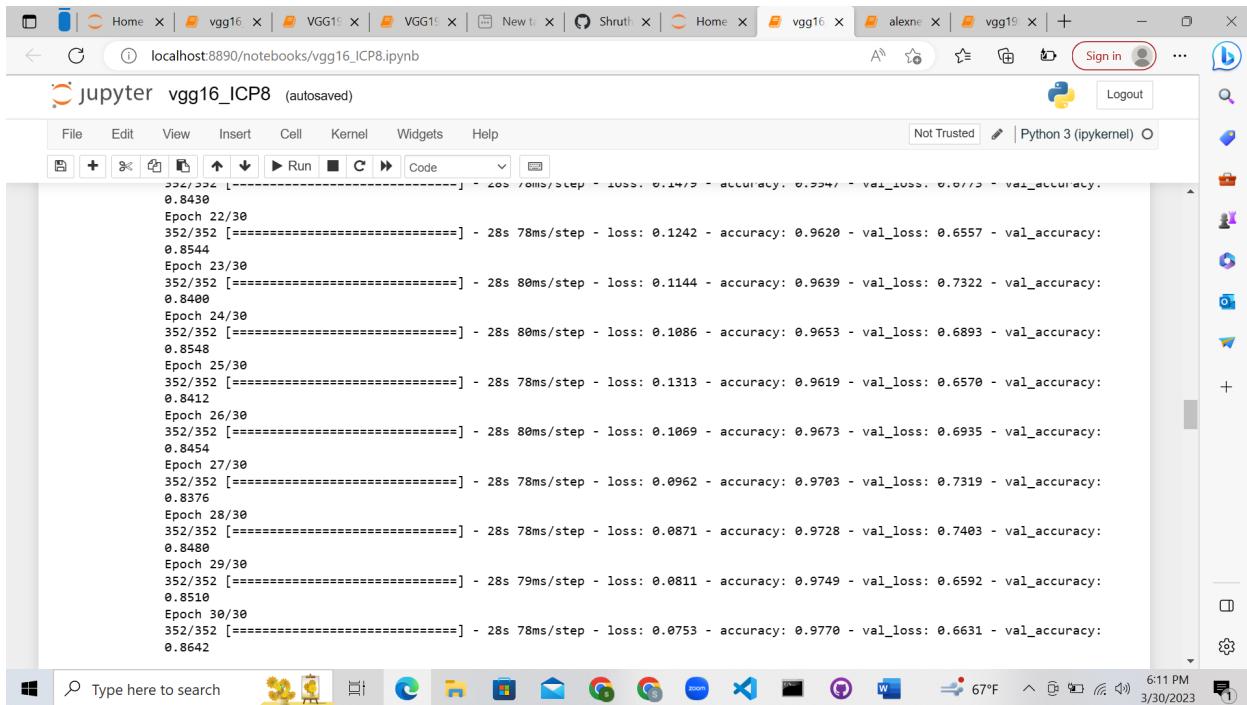
Total params: 27,331,914
Trainable params: 27,323,466
Non-trainable params: 8,448

In []: `history = model.fit(x=x_train, y=one_hot_y_train, batch_size=128, epochs=30, validation_split=0.1)`

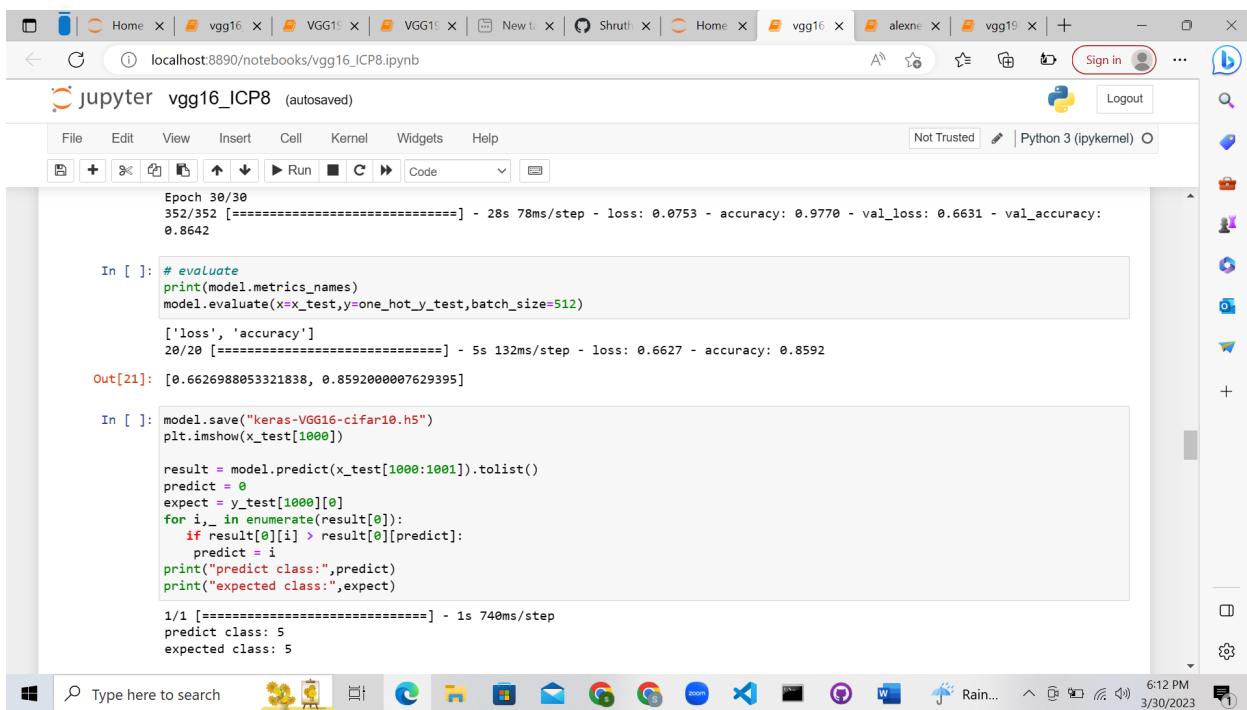


In []: `history = model.fit(x=x_train, y=one_hot_y_train, batch_size=128, epochs=30, validation_split=0.1)`

```
Epoch 1/30
352/352 [=====] - 29s 81ms/step - loss: 1.7083 - accuracy: 0.3346 - val_loss: 2.6400 - val_accuracy: 0.3686
Epoch 2/30
352/352 [=====] - 27s 76ms/step - loss: 1.3159 - accuracy: 0.5254 - val_loss: 1.6935 - val_accuracy: 0.5590
Epoch 3/30
352/352 [=====] - 27s 77ms/step - loss: 1.0110 - accuracy: 0.6538 - val_loss: 1.0628 - val_accuracy: 0.6556
Epoch 4/30
352/352 [=====] - 27s 78ms/step - loss: 0.8493 - accuracy: 0.7144 - val_loss: 0.9995 - val_accuracy: 0.6856
Epoch 5/30
352/352 [=====] - 28s 79ms/step - loss: 0.7343 - accuracy: 0.7566 - val_loss: 0.8316 - val_accuracy: 0.7326
Epoch 6/30
352/352 [=====] - 28s 79ms/step - loss: 0.6515 - accuracy: 0.7907 - val_loss: 0.8145 - val_accuracy: 0.7648
Epoch 7/30
352/352 [=====] - 28s 80ms/step - loss: 0.5708 - accuracy: 0.8168 - val_loss: 0.7658 - val_accuracy: 0.7716
Epoch 8/30
352/352 [=====] - 28s 79ms/step - loss: 0.5135 - accuracy: 0.8339 - val_loss: 0.6754 - val_accuracy: 0.7944
Epoch 9/30
352/352 [=====] - 28s 79ms/step - loss: 0.4665 - accuracy: 0.8506 - val_loss: 0.7615 - val_accuracy: 0.7686
```



```
0.8430
Epoch 22/30
352/352 [=====] - 28s 78ms/step - loss: 0.1242 - accuracy: 0.9620 - val_loss: 0.6557 - val_accuracy: 0.8544
Epoch 23/30
352/352 [=====] - 28s 80ms/step - loss: 0.1144 - accuracy: 0.9639 - val_loss: 0.7322 - val_accuracy: 0.8400
Epoch 24/30
352/352 [=====] - 28s 80ms/step - loss: 0.1086 - accuracy: 0.9653 - val_loss: 0.6893 - val_accuracy: 0.8544
Epoch 25/30
352/352 [=====] - 28s 78ms/step - loss: 0.1313 - accuracy: 0.9619 - val_loss: 0.6570 - val_accuracy: 0.8412
Epoch 26/30
352/352 [=====] - 28s 80ms/step - loss: 0.1069 - accuracy: 0.9673 - val_loss: 0.6935 - val_accuracy: 0.8454
Epoch 27/30
352/352 [=====] - 28s 78ms/step - loss: 0.0962 - accuracy: 0.9703 - val_loss: 0.7319 - val_accuracy: 0.8376
Epoch 28/30
352/352 [=====] - 28s 78ms/step - loss: 0.0871 - accuracy: 0.9728 - val_loss: 0.7403 - val_accuracy: 0.8480
Epoch 29/30
352/352 [=====] - 28s 79ms/step - loss: 0.0811 - accuracy: 0.9749 - val_loss: 0.6592 - val_accuracy: 0.8510
Epoch 30/30
352/352 [=====] - 28s 78ms/step - loss: 0.0753 - accuracy: 0.9770 - val_loss: 0.6631 - val_accuracy: 0.8642
```



```
In [ ]: # evaluate
print(model.metrics_names)
model.evaluate(x=x_test,y=one_hot_y_test,batch_size=512)
['loss', 'accuracy']
20/20 [=====] - 5s 132ms/step - loss: 0.6627 - accuracy: 0.8592

Out[21]: [0.6626988053321838, 0.8592000007629395]

In [ ]: model.save("keras-VGG16-cifar10.h5")
plt.imshow(x_test[1000])

result = model.predict(x_test[1000:1001]).tolist()
predict = 0
expect = y_test[1000][0]
for i,_ in enumerate(result[0]):
    if result[0][i] > result[0][predict]:
        predict = i
print("predict class:",predict)
print("expected class:",expect)

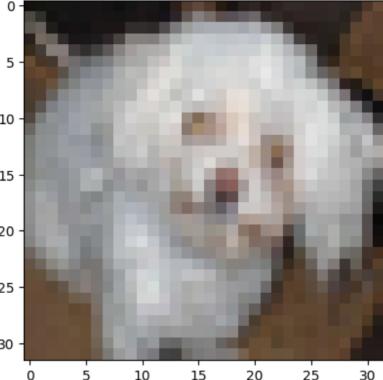
1/1 [=====] - 1s 740ms/step
predict class: 5
expected class: 5
```

jupyter vgg16_ICP8 (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted Python 3 (ipykernel) ○

```
1/1 [=====] - 740ms/step
predict class: 5
expected class: 5
```



6:12 PM 3/30/2023

jupyter vgg16_ICP8 (autosaved)

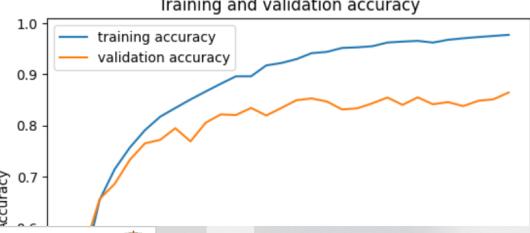
File Edit View Insert Cell Kernel Widgets Help

Not Trusted Python 3 (ipykernel) ○

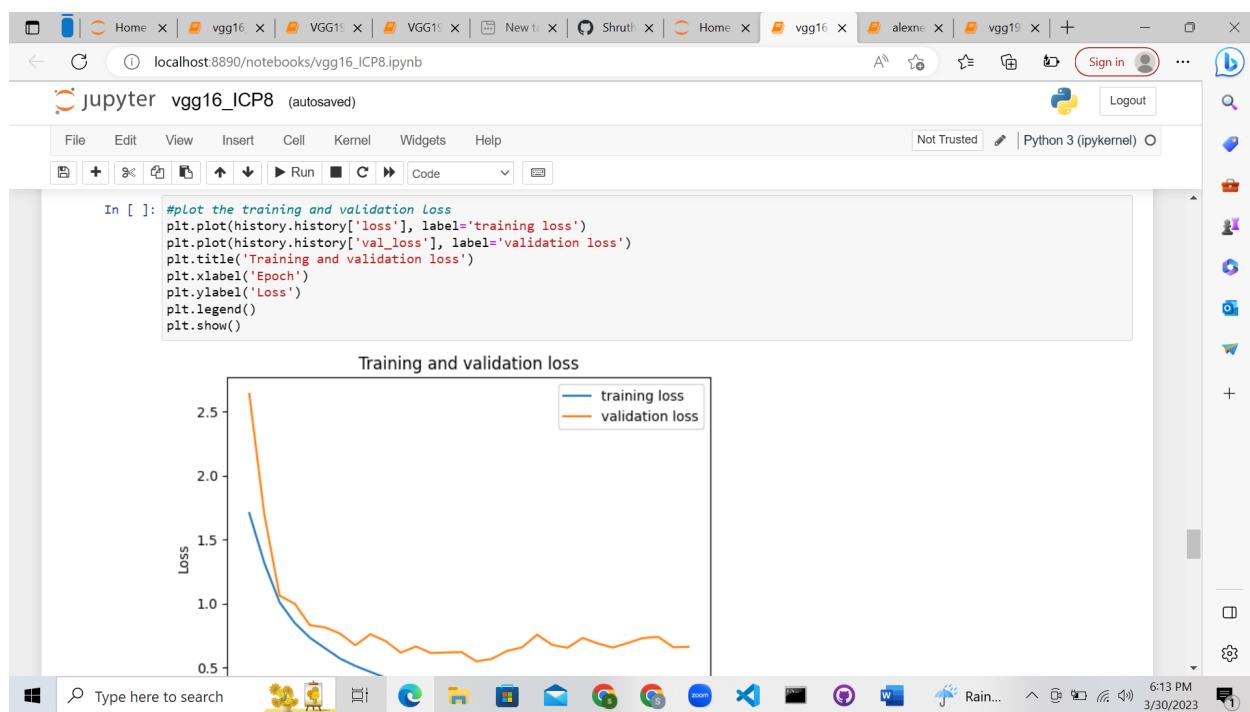
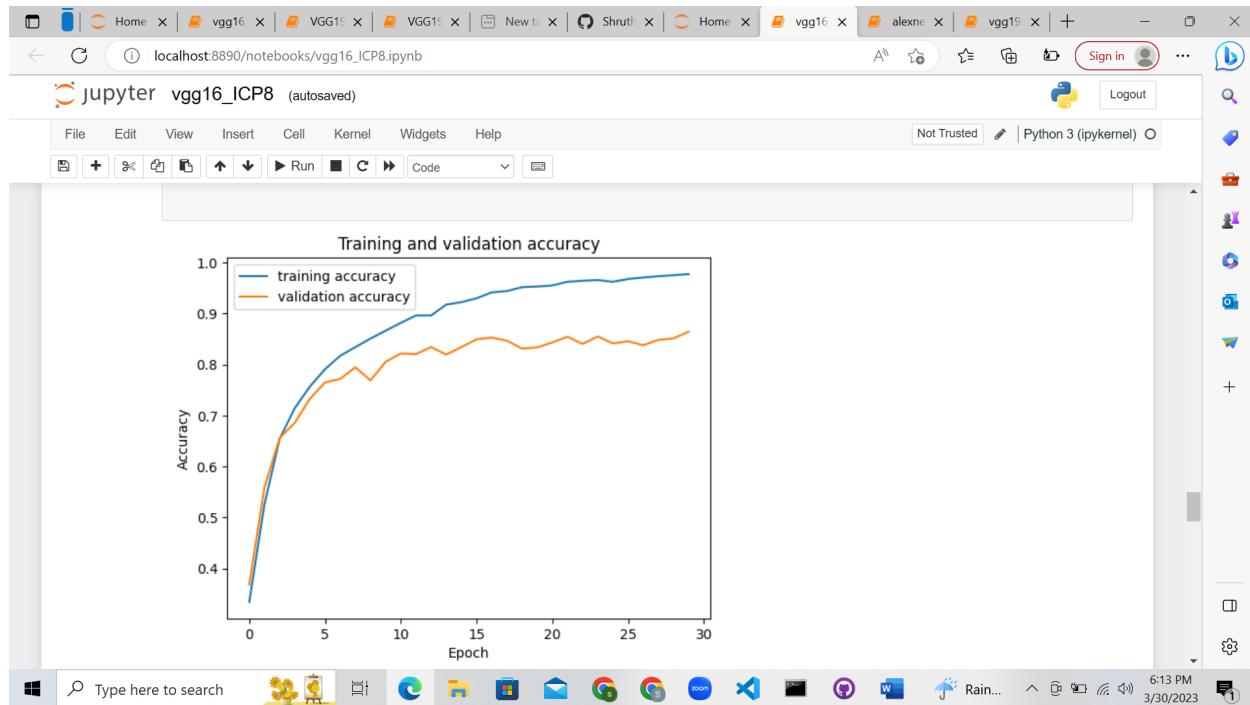
```
In [16]: # save model
model.save("keras-VGG16-cifar10.h5")
```

```
In [ ]: #plot the training and validation accuracy
plt.plot(history.history['accuracy'], label='training accuracy')
plt.plot(history.history['val_accuracy'], label='validation accuracy')
plt.title('Training and validation accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

Training and validation accuracy



6:13 PM 3/30/2023



The screenshot shows a Jupyter Notebook interface running on a Windows desktop. The notebook is titled "jupyter vgg16_ICP8 (autosaved)". The code cell contains Python code for calculating a confusion matrix and plotting predicted probabilities. The code uses NumPy and the sklearn.metrics module. It includes comments explaining the steps: calculating the confusion matrix, plotting it with a colorbar, and plotting a histogram of predicted probabilities for a sample image. The status bar at the bottom shows "313/313 [=====] - 3s 8ms/step". The taskbar at the bottom of the screen shows various pinned icons, including a search bar, a calendar, and several Microsoft applications like Edge, Mail, and File Explorer.

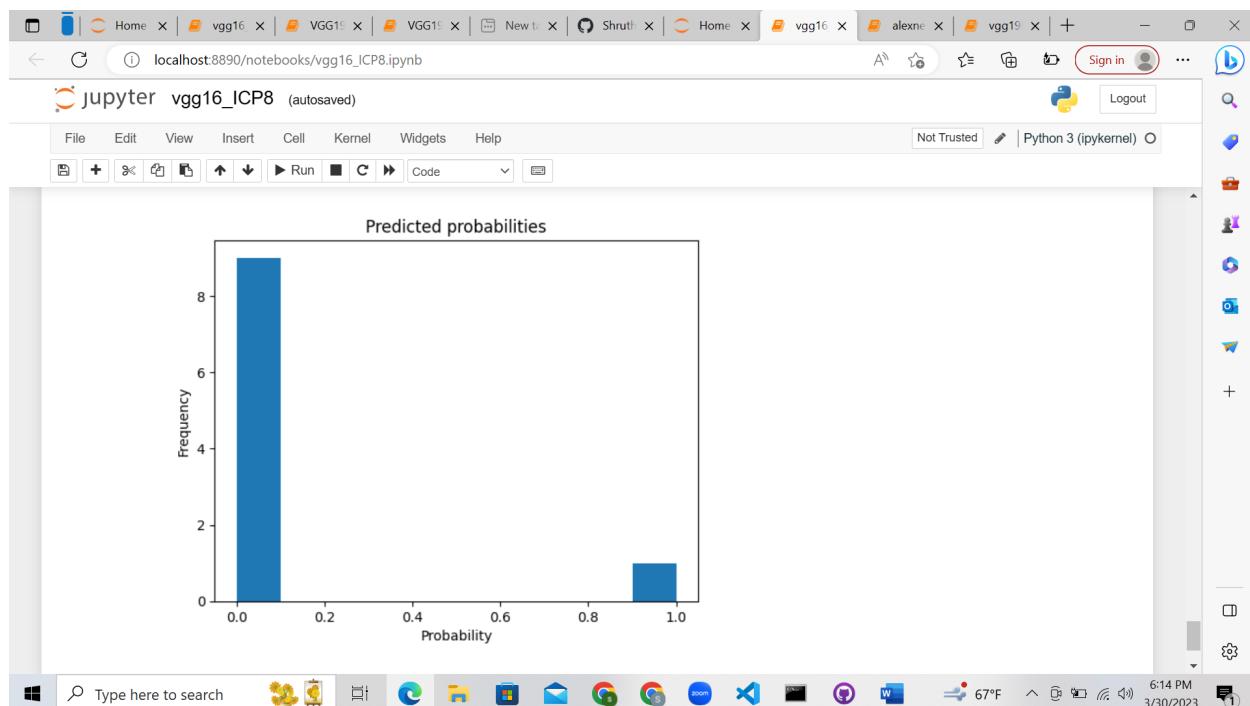
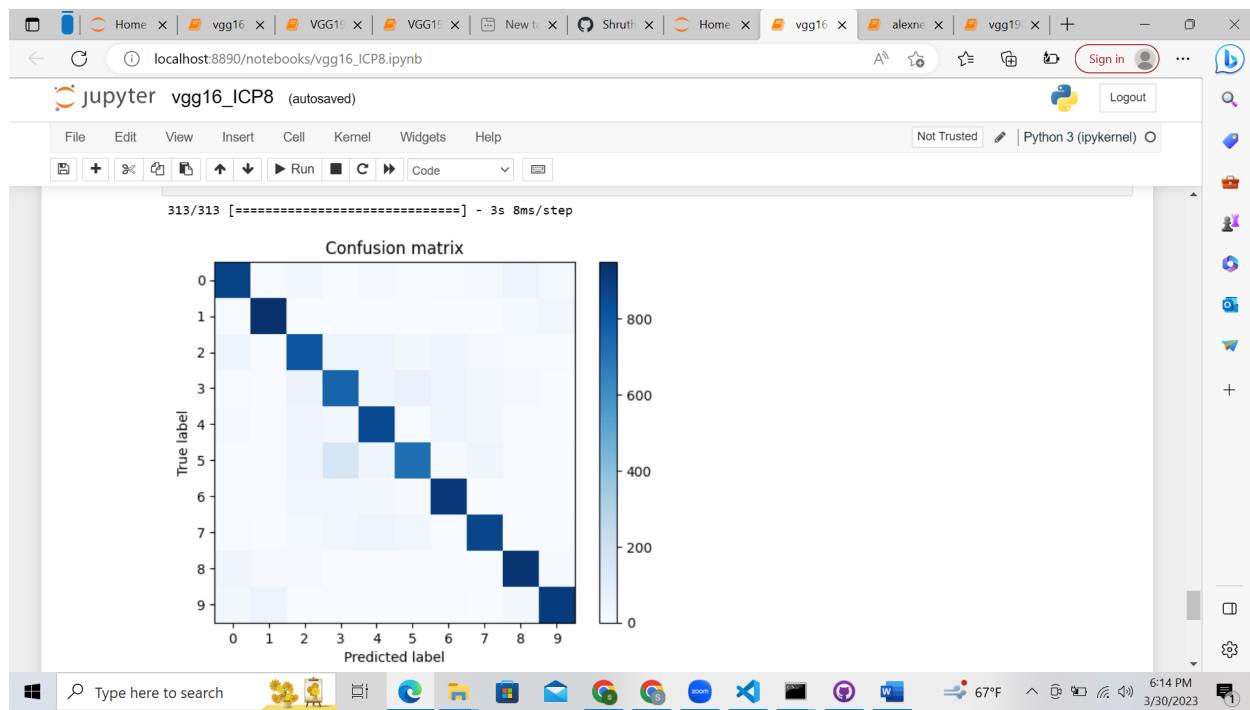
```
In [ ]: import numpy as np
from sklearn.metrics import confusion_matrix

# calculate the confusion matrix
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = y_test.ravel()
cm = confusion_matrix(y_true, y_pred_classes)

# plot the confusion matrix
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion matrix')
plt.colorbar()
tick_marks = np.arange(num_classes)
plt.xticks(tick_marks, range(num_classes))
plt.yticks(tick_marks, range(num_classes))
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.show()

# plot a histogram of the predicted probabilities for a sample image
plt.hist(y_pred[1000])
plt.title('Predicted probabilities')
plt.xlabel('Probability')
plt.ylabel('Frequency')
plt.show()
```

313/313 [=====] - 3s 8ms/step



Model after improvising

localhost:8888/notebooks/VGG19Model_(CIFAR100)%20(4).ipynb

jupyter VGG19Model_(CIFAR100) (4) Last Checkpoint: 6 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted | Python 3 (ipykernel) ○

Import libraries

```
In [11]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow.keras.datasets import mnist

from tensorflow.keras.optimizers import RMSprop
from keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalization

%matplotlib inline
```

Extract data and train and test dataset

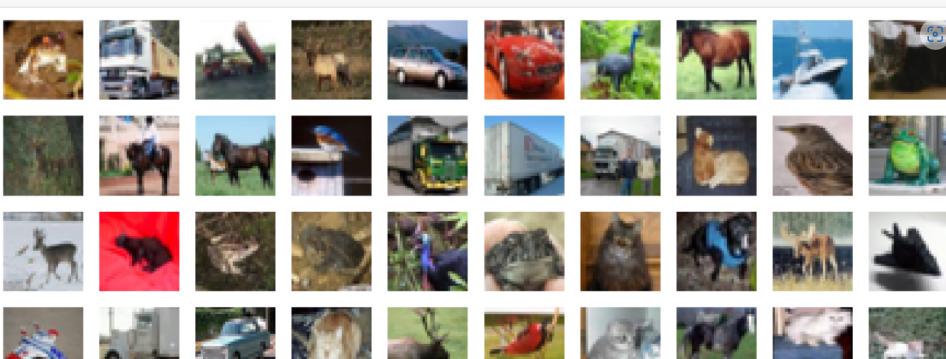
```
In [12]: #cifar100 = tf.keras.datasets.cifar100
(X_train,Y_train) , (X_test,Y_test) = cifar100.load_data()

In [13]: classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

Let's look into the dataset images

Type here to search

```
In [14]: plt.figure(figsize = (16,16))
for i in range(100):
    plt.subplot(10,10,i+1)
    plt.axis('off')
    plt.imshow(X_train[i], cmap = 'gray')
```



6:15 PM 3/30/2023

localhost:8888/notebooks/VGG19Model_(CIFAR100)%20(4).ipynb

jupyter VGG19Model_(CIFAR100) (4) Last Checkpoint: 6 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted | Python 3 (ipykernel) ○

localhost:8888/notebooks/VGG19Model_(CIFAR100)%20(4).ipynb

jupyter VGG19Model_(CIFAR100) (4) Last Checkpoint: 6 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted | Python 3 (ipykernel) O

Training , Validating and Splitting trained and tested data

```
In [15]: from sklearn.model_selection import train_test_split
x_train, x_val, y_train, y_val = train_test_split(X_train,Y_train,test_size=0.2)

In [16]: from keras.utils.np_utils import to_categorical
y_train = to_categorical(y_train, num_classes = 10)
y_val = to_categorical(y_val, num_classes = 10)

In [17]: print(x_train.shape)
print(y_train.shape)
print(x_val.shape)
print(y_val.shape)
print(X_test.shape)
print(Y_test.shape)

(40000, 32, 32, 3)
(40000, 10)
(10000, 32, 32, 3)
(10000, 10)
(10000, 32, 32, 3)
(10000, 1)

In [18]: train_datagen = ImageDataGenerator(
    preprocessing_function = tf.keras.applications.vgg19.preprocess_input,
    rotation_range=10,
    zoom_range = 0.1,
```

6:16 PM 3/30/2023

localhost:8888/notebooks/VGG19Model_(CIFAR100)%20(4).ipynb

jupyter VGG19Model_(CIFAR100) (4) Last Checkpoint: 6 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted | Python 3 (ipykernel) ○

```
In [18]: train_datagen = ImageDataGenerator(  
    preprocessing_function = tf.keras.applications.vgg19.preprocess_input,  
    rotation_range=10,  
    zoom_range = 0.1,  
    width_shift_range = 0.1,  
    height_shift_range = 0.1,  
    shear_range = 0.1,  
    horizontal_flip = True  
)  
train_datagen.fit(x_train)  
  
val_datagen = ImageDataGenerator(preprocessing_function = tf.keras.applications.vgg19.preprocess_input)  
val_datagen.fit(x_val)
```

In [19]: from keras.callbacks import ReduceLROnPlateau
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
 patience=3,
 verbose=1,
 factor=0.5,
 min_lr=0.00001)

We have used only 16 layers out of 19 layers in the CNN

In [20]: vgg_model = tf.keras.applications.VGG19(
 include_top=False,
 weights=None,
 input_shape=(32,32,3))

localhost:8888/notebooks/VGG19Model_(CIFAR100)%20(4).ipynb

jupyter VGG19Model_(CIFAR100) (4) Last Checkpoint: 6 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted | Python 3 (ipykernel) ○

```
In [20]: vgg_model = tf.keras.applications.VGG19(  
    include_top=False,  
    weights=None,  
    input_shape=(32,32,3),  
)  
  
vgg_model.summary()
```

Model: "vgg19"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080

localhost:8888/notebooks/VGG19Model_(CIFAR100)%20(4).ipynb

jupyter VGG19Model_(CIFAR100) (4) Last Checkpoint: 6 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted | Python 3 (ipykernel) ○

Run Cell Cell Kernel Help

Layer	Type	Shape	Param #
block4_conv1	Conv2D	(None, 4, 4, 512)	1180160
block4_conv2	Conv2D	(None, 4, 4, 512)	2359808
block4_conv3	Conv2D	(None, 4, 4, 512)	2359808
block4_conv4	Conv2D	(None, 4, 4, 512)	2359808
block4_pool	MaxPooling2D	(None, 2, 2, 512)	0
block5_conv1	Conv2D	(None, 2, 2, 512)	2359808
block5_conv2	Conv2D	(None, 2, 2, 512)	2359808
block5_conv3	Conv2D	(None, 2, 2, 512)	2359808
block5_conv4	Conv2D	(None, 2, 2, 512)	2359808
block5_pool	MaxPooling2D	(None, 1, 1, 512)	0
<hr/>			
Total params: 20,024,384			
Trainable params: 20,024,384			
Non-trainable params: 0			

In [24]: model = tf.keras.Sequential()
model.add(vgg_model)

6:17 PM 3/30/2023

localhost:8888/notebooks/VGG19Model_(CIFAR100)%20(4).ipynb

jupyter VGG19Model_(CIFAR100) (4) Last Checkpoint: 6 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted | Python 3 (ipykernel) ○

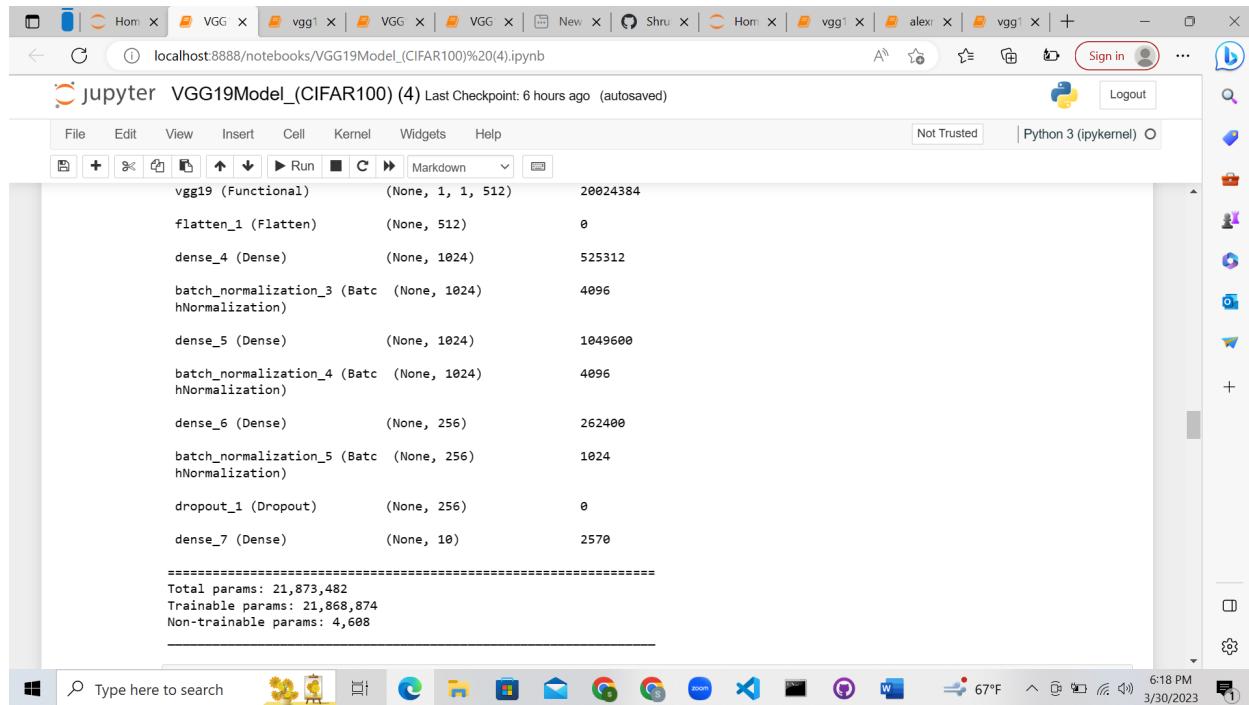
Run Cell Cell Kernel Help

```
In [24]: model = tf.keras.Sequential()  
model.add(vgg_model)  
model.add(Flatten())  
model.add(Dense(1024, activation = 'relu'))  
model.add(BatchNormalization())  
model.add(Dense(1024, activation = 'relu'))  
model.add(BatchNormalization())  
model.add(Dense(256, activation = 'relu'))  
model.add(BatchNormalization())  
model.add(Dropout(0.5))  
model.add(Dense(10, activation = 'softmax'))  
  
model.summary()
```

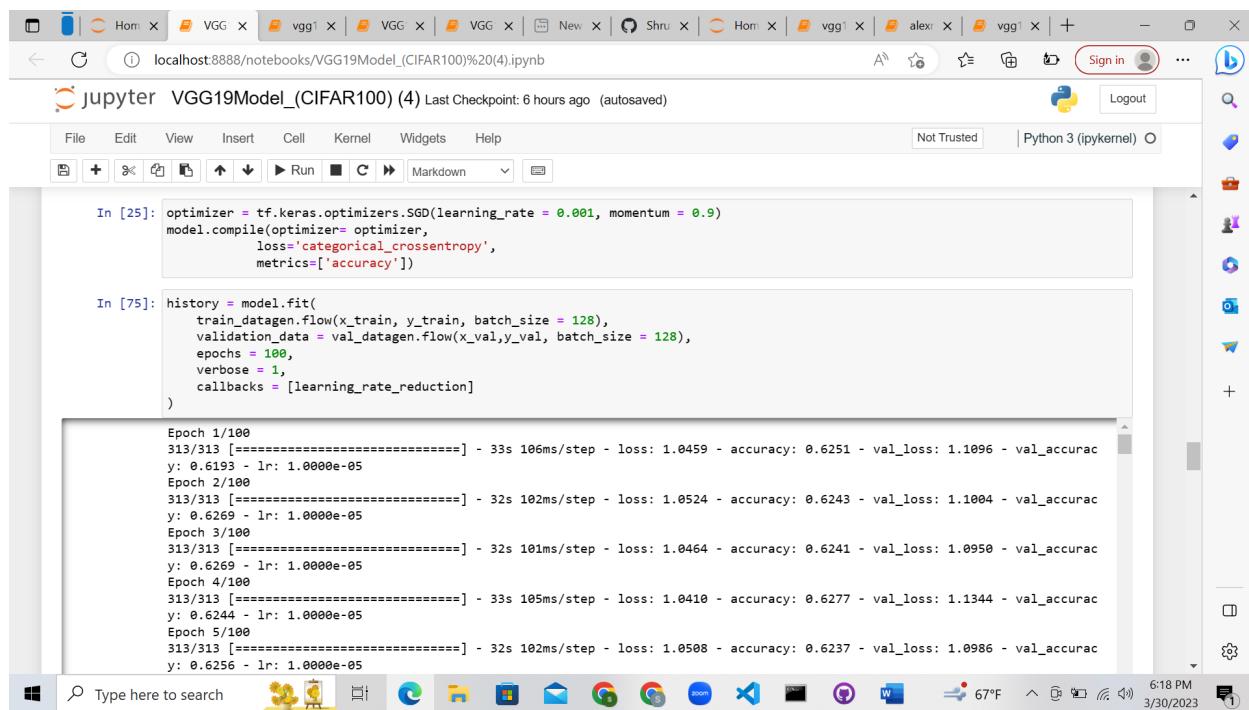
Model: "sequential_1"

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 1, 1, 512)	20024384
flatten_1 (Flatten)	(None, 512)	0
dense_4 (Dense)	(None, 1024)	525312
batch_normalization_3 (BatchNormalization)	(None, 1024)	4096
dense_5 (Dense)	(None, 1024)	1049600

6:17 PM 3/30/2023



```
vgg19 (Functional)      (None, 1, 512)      20024384
flatten_1 (Flatten)     (None, 512)          0
dense_4 (Dense)         (None, 1024)         525312
batch_normalization_3 (Batch Normalization) (None, 1024) 4096
dense_5 (Dense)         (None, 1024)         1049600
batch_normalization_4 (Batch Normalization) (None, 1024) 4096
dense_6 (Dense)         (None, 256)          262400
batch_normalization_5 (Batch Normalization) (None, 256) 1024
dropout_1 (Dropout)     (None, 256)          0
dense_7 (Dense)         (None, 10)           2570
=====
Total params: 21,873,482
Trainable params: 21,868,874
Non-trainable params: 4,608
```



```
In [25]: optimizer = tf.keras.optimizers.SGD(learning_rate = 0.001, momentum = 0.9)
model.compile(optimizer=optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])

In [75]: history = model.fit(
            train_datagen.flow(x_train, y_train, batch_size = 128),
            validation_data = val_datagen.flow(x_val,y_val, batch_size = 128),
            epochs = 100,
            verbose = 1,
            callbacks = [learning_rate_reduction])
```

```
Epoch 1/100
313/313 [=====] - 33s 106ms/step - loss: 1.0459 - accuracy: 0.6251 - val_loss: 1.1096 - val_accuracy: 0.6193 - lr: 1.0000e-05
Epoch 2/100
313/313 [=====] - 32s 102ms/step - loss: 1.0524 - accuracy: 0.6243 - val_loss: 1.1004 - val_accuracy: 0.6269 - lr: 1.0000e-05
Epoch 3/100
313/313 [=====] - 32s 101ms/step - loss: 1.0464 - accuracy: 0.6241 - val_loss: 1.0950 - val_accuracy: 0.6269 - lr: 1.0000e-05
Epoch 4/100
313/313 [=====] - 33s 105ms/step - loss: 1.0410 - accuracy: 0.6277 - val_loss: 1.1344 - val_accuracy: 0.6244 - lr: 1.0000e-05
Epoch 5/100
313/313 [=====] - 32s 102ms/step - loss: 1.0508 - accuracy: 0.6237 - val_loss: 1.0986 - val_accuracy: 0.6256 - lr: 1.0000e-05
```

localhost:8888/notebooks/VGG19Model_(CIFAR100)%20(4).ipynb

jupyter VGG19Model_(CIFAR100) (4) Last Checkpoint: 6 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted Python 3 (ipykernel) O

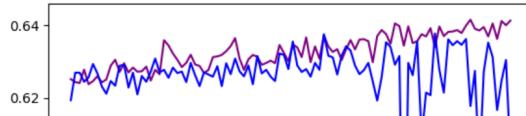
```
313/313 [=====] - 32s 102ms/step - loss: 1.0142 - accuracy: 0.6362 - val_loss: 1.1493 - val_accuracy: 0.6166 - lr: 1.0000e-05
Epoch 98/100
313/313 [=====] - 32s 103ms/step - loss: 1.0011 - accuracy: 0.6411 - val_loss: 1.1204 - val_accuracy: 0.6247 - lr: 1.0000e-05
Epoch 99/100
313/313 [=====] - 33s 105ms/step - loss: 1.0009 - accuracy: 0.6400 - val_loss: 1.0980 - val_accuracy: 0.6304 - lr: 1.0000e-05
Epoch 100/100
313/313 [=====] - 33s 104ms/step - loss: 0.9992 - accuracy: 0.6413 - val_loss: 1.1563 - val_accuracy: 0.6038 - lr: 1.0000e-05
```

In [76]:

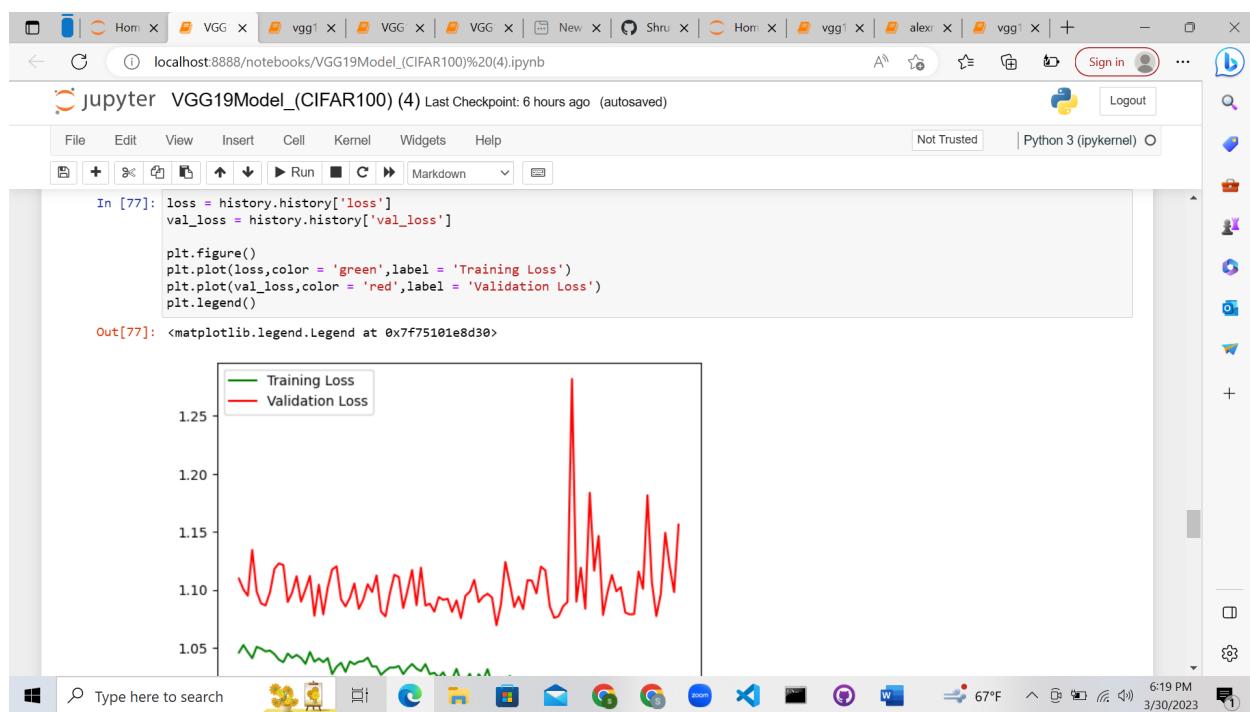
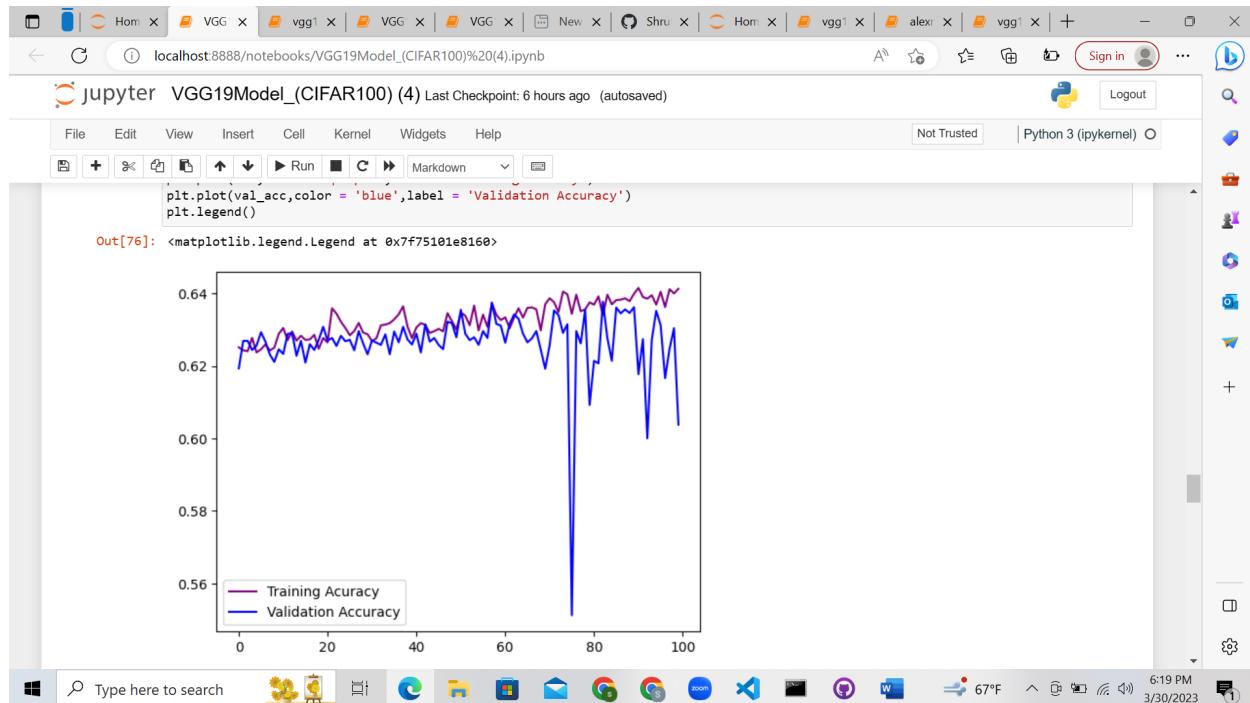
```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

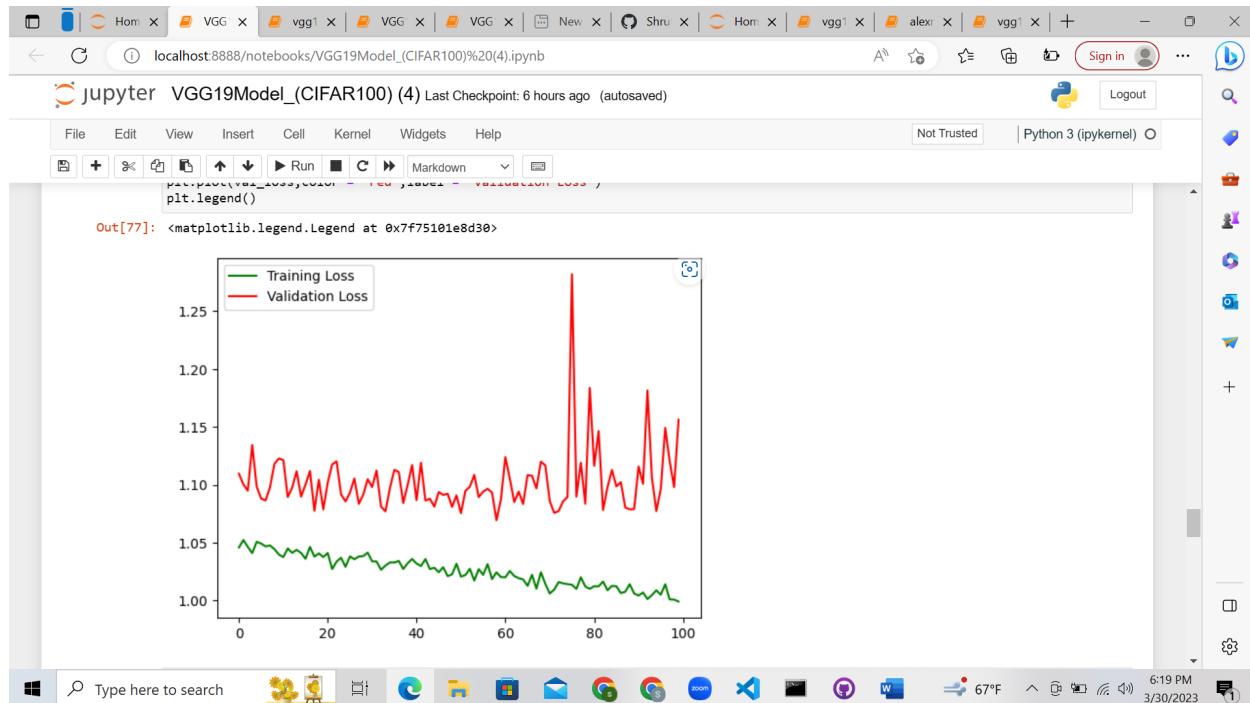
plt.figure()
plt.plot(acc,color = 'purple',label = 'Training Accuracy')
plt.plot(val_acc,color = 'blue',label = 'Validation Accuracy')
plt.legend()
```

Out[76]:



67°F 6:18 PM 3/30/2023





```
In [78]: X_test = tf.keras.applications.vgg19.preprocess_input(X_test)
y_pred = np.argmax(model.predict(X_test), axis=-1)

y_pred[:10]
313/313 [=====] - 3s 9ms/step

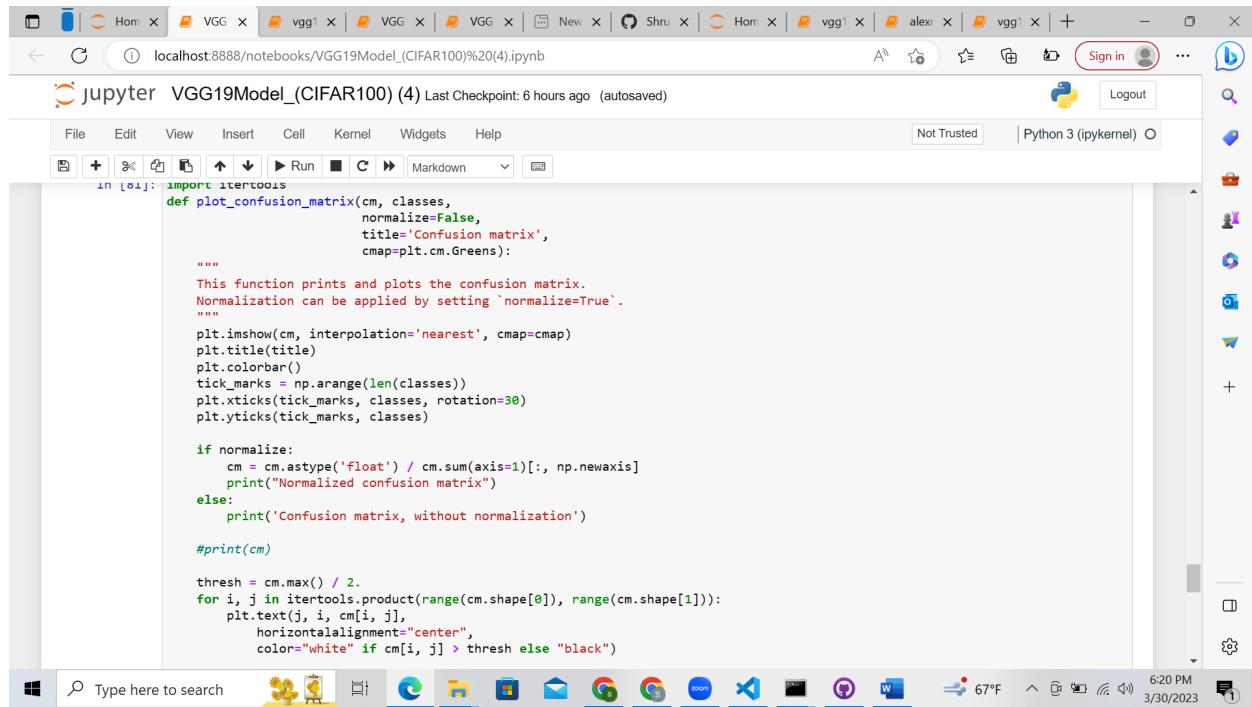
Out[78]: array([5, 1, 5, 5, 5, 5, 7, 5, 5, 7])

In [79]: from sklearn.metrics import confusion_matrix, accuracy_score
print('Testing Accuracy : ', accuracy_score(Y_test, y_pred))

Testing Accuracy : 0.1326

In [80]: cm = confusion_matrix(Y_test, y_pred)
cm
```

```
Out[80]: array([[ 7, 36,  0,  4,  0, 433,  0, 484,  0, 36],
   [ 61, 141,  0,  2,  1, 250,  0, 399,  0, 146],
   [ 1,  0,  0, 10,  0, 737,  0, 250,  0, 2],
   [ 0,  1,  0,  8,  3, 685,  0, 295,  0, 8],
   [ 0,  0,  16,  3, 779,  0, 197,  0, 5],
   [ 1,  0,  0, 19,  2, 684,  0, 290,  0, 4],
   [ 3,  5,  0,  9,  1, 716,  0, 252,  0, 14],
   [ 1,  5,  0, 18,  2, 597,  0, 334,  0, 43],
   [ 15, 34,  0,  3,  0, 469,  0, 423,  0, 56],
   [ 27, 149,  0,  2,  1, 319,  0, 353,  0, 149]])
```



The screenshot shows a Jupyter Notebook interface running on a Windows desktop. The title bar indicates the notebook is titled "VGG19Model_(CIFAR100) (4)" and was last checked 6 hours ago. The status bar at the bottom shows the date as 3/30/2023 and the time as 6:20 PM. The main content area displays a Python script for plotting a confusion matrix. The script imports `itertools` and defines a function `plot_confusion_matrix` that takes a confusion matrix `cm` and optional parameters `normalize`, `title`, and `cmap`. The function prints the matrix and plots it using `plt.imshow`. It includes a docstring explaining the function's purpose and normalization options. The script then defines a threshold and iterates over the matrix to print values above the diagonal. The code is as follows:

```
in [81]: import itertools
def plot_confusion_matrix(cm, classes,
                         normalize=False,
                         title='Confusion matrix',
                         cmap=plt.cm.Greens):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=30)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    #print(cm)

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")
```

jupyter VGG19Model_(CIFAR100) (4) Last Checkpoint: 6 hours ago (autosaved)

```
File Edit View Insert Cell Kernel Widgets Help
Not Trusted Python 3 (ipykernel) ○
```

```
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=30)
plt.yticks(tick_marks, classes)

if normalize:
    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
    print("Normalized confusion matrix")
else:
    print('Confusion matrix, without normalization')

#print(cm)

thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, cm[i, j],
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

In [82]: plt.figure(figsize=(8,8))
plot_confusion_matrix(cm,classes)
```

Confusion matrix, without normalization

